

aslbench Run Report

Table of contents

Run metadata	1
Comparison	1
McNemar test	4
Per-model detail	5
gpt-5.6-sol	5
gpt-5.6-terra	6
claude-opus-4.8	6
claude-sonnet-4.6	7

Run metadata

- **Run:** aslbench-20260713-203909
- **Started:** 2026-07-13T20:39:09
- **Images per class:** 10
- **Classes:** 36
- **Total images:** 360
- **Sample seed:** 156084346
- **Prompt template:** v2_class_list
- **Note:** GPT-5.6 Sol, GPT-5.6 Terra, Opus 4.8, Sonnet 4.6


Provider	Model	Slug
GitHub Copilot	gpt-5.6-sol	copilot-gpt-5-6-sol
GitHub Copilot	gpt-5.6-terra	copilot-gpt-5-6-terra
GitHub Copilot	claude-opus-4.8	copilot-claude-opus-4-8
GitHub Copilot	claude-sonnet-4.6	copilot-claude-sonnet-4-6

The task is 36-way single-character classification (ASL fingerspelling digits 0-9 and letters A-Z). Every model saw the identical sampled images, so all differences are attributable to the model rather than the sample. Filenames, which encode the true class, are never shown to the models.

Comparison


metric	gpt-5.6-sol	gpt-5.6-terra	claude-opus-4.8	claude-sonnet-4.6	Chance (1/36)
accuracy	0.561	0.453	0.397	0.317	0.028
macro_f1	0.517	0.386	0.355	0.249	0.028
parse_failure_rate	0.000	0.000	0.000	0.000	0.000
provider_error_rate	0.000	0.000	0.000	0.000	0.000

- **Accuracy** — the share of images the model labelled correctly, from 0 to 1; simply how often it is right.
- **Macro F1** — the model’s reliability averaged across all 36 characters, giving each character equal weight and blending how often its guesses are right (precision) with how often it finds each character (recall). It rewards models that do well on every character, not just the easy ones.
- **Parse failure rate** — how often the model’s reply could not be read as a valid answer (it did not clearly name one of the 36 characters).
- **Provider error rate** — how often the model’s API failed to return any response at all (timeouts or errors).
- **Chance (1/36)** — a reference column showing what each score would be for a model that guesses at random.




Per-class accuracy difference (sorted by advantage)

Model ■ gpt-5.6-sol better ■ gpt-5.6-terra better ■ claude-opus-4.8 better ■ claude-sonnet-4.6 better



Confusion matrix (row-normalized, recall)



How to read this: each row is the true character and each column is what the model guessed. Cells are row-normalized, so a value is the fraction of that character's images that received the correct guess; the diagonal is the model's recall for each character. A bright diagonal means accurate recognition, while bright off-diagonal cells show characters the model routinely mixes up (e.g. reading 'O' as '0').

McNemar test

Model A	Model B	Better	A only	B only	Discor- dant	p	p (Holm)	Sig.
gpt-5.6-sol	claude-son- net-4.6	gpt-5.6-sol	105	17	122	0	0	yes
gpt-5.6-sol	claude- opus-4.8	gpt-5.6-sol	87	28	115	0	0	yes
gpt-5.6- terra	claude-son- net-4.6	gpt-5.6- terra	78	29	107	0	0	yes
gpt-5.6-sol	gpt-5.6-terra	gpt-5.6-sol	67	28	95	0	0	yes
claude- opus-4.8	claude-son- net-4.6	claude- opus-4.8	61	32	93	0.003	0.007	yes
gpt-5.6- terra	claude- opus-4.8	gpt-5.6- terra	55	35	90	0.045	0.045	yes

What this answers: whether two models have genuinely different accuracy, judged only on the images where they disagree. Because every model saw the exact same images, we compare them image-by-image. The test looks only at *discordant items*, that is, images where one model was right and the other was wrong. Images they both got right or both got wrong say nothing about which is better, so they are set aside. If those disagreements are far more lopsided toward one model than a coin flip would explain, the difference is judged real. In the table, **better** names the model ahead on the discordant images, **only_a_correct** and **only_b_correct** count how many each won, **p_value** is the raw test result, and **p_holm** is that p-value after a Holm-Bonferroni correction that accounts for comparing every pair of models at once. **significant** is yes when **p_holm** is below 0.05. Rows are sorted most-significant first.

Per-model detail

gpt-5.6-sol

Most-confused pairs

true_char	pred_char	count
J	I	10
2	V	9
M	S	9
N	S	8
K	V	7
E	S	7
7	3	6
8	7	5

Sample misclassifications

item_id	true_char	predicted_char
P2_0_174	0	6
P5_0_426	0	6
P7_0_654	0	6
P8_0_751	0	B
P8_0_771	0	B
P9_0_865	0	6
P5_1_483	1	D
P5_1_485	1	D

gpt-5.6-terra**Most-confused pairs**

true_char	pred_char	count
M	S	10
1	D	10
J	I	10
F	W	9
2	V	9
P	G	9
N	S	9
9	W	8

Sample misclassifications

item_id	true_char	predicted_char
P3_0_239	0	6
P8_0_771	0	4
P10_1_928	1	D
P10_1_969	1	D
P2_1_138	1	D
P2_1_143	1	D
P5_1_483	1	D
P5_1_485	1	D

claude-opus-4.8

Most-confused pairs

true_char	pred_char	count
2	V	10
K	V	9
N	S	8
1	D	7
E	S	7
Q	X	6
T	S	6
M	S	5

Sample misclassifications

item_id	true_char	predicted_char
P2_0_174	0	7
P3_0_239	0	F
P8_0_751	0	1
P8_0_771	0	6
P9_0_845	0	B
P9_0_865	0	6
P10_1_928	1	L
P10_1_969	1	D

claude-sonnet-4.6**Most-confused pairs**

true_char	pred_char	count
K	V	10
2	V	10
J	Y	8
E	S	8
D	1	8
S	A	6
3	V	6
P	G	6

Sample misclassifications

item_id	true_char	predicted_char
P1_0_21	0	9
P1_0_30	0	F
P2_0_174	0	B
P3_0_239	0	U
P5_0_426	0	F
P7_0_654	0	V
P8_0_751	0	B
P8_0_771	0	H